

End-to-end inference ~ 20 ms (~ 50 FPS) on a Quadro RTX 8000

